




Part 4

Arithmetic Logic Unit




Negative Binary Integers

Have a positive attitude about negatives

Negative Binary Numbers

- When we write a negative number, we generally use a "-" as a prefix character
- However, binary numbers can only store ones and zeros




6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

3

Negative Binary Numbers

- So, how we store a negative a number?
- When a number can represent both positive and negative numbers, it is called a *signed integer*
- Otherwise, it is *unsigned*



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

4

Signed Magnitude


- One approach is to use the most significant bit (msb) to represent the negative sign
- If positive, this bit will be a zero
- If negative, this bit will be a 1
- This gives a byte a range of -127 to 127 rather than 0 to 255

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

5

Signed Magnitude




6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019


6

Signed Magnitude: 13 and -13

Positive



Negative



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

7

Signed Magnitude Drawback #1

- When two numbers are added, the system needs to check and sign bits and act accordingly
- For example:
 - if both numbers are positive, add values
 - if one is negative subtract it from the other
 - etc...
- There are also rules for subtracting

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

8

Signed Magnitude Drawback #2

- Also, signed magnitude also can store a positive *and* negative version of zero
- Yes, there are two zeroes!
- Imagine having to write Java code like...

```
if (x == +0 || x == -0)
```


6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019


9

Oh noes! Two zeros?

+0



-0



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

10

1's Complement

- Rather than use a sign bit, the value can be made negative by *inverting* each bit
 - each 1 becomes a 0
 - each 0 becomes a 1
- Result is a "complement" of the original
- This is logically the same as subtracting the number from 0

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

11

Advantages / Disadvantages

- Advantages over signed magnitude
 - very simple rules for adding/subtracting
 - numbers are simply added: 5 - 3 is the same as 5 + -3
- Disadvantages
 - positive and negative zeros still exist
 - so, it's not a perfect solution

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

12

1's Complement: 13 and -13

Positive

0 0 0 0 1 1 0 1

Negative

1 1 1 1 0 0 1 0

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

13

1's Complement Has Two Zeros

+0

0 0 0 0 0 0 0 0

-0

1 1 1 1 1 1 1 1

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

14

2's Complement

- Practically all computers nowadays use *2's Complement*
- Similar to 1's complement, but after the number is inverted, 1 is added to the result
- Logically the same as:
 - subtracting the number from 2^n
 - where n is the total number of bits in the integer

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

15

2's Complement Advantages

- Since negatives are subtracted from 2^n
 - they can simply be added
 - the extra carry 1 (if it exists) is discarded
 - this simplifies the hardware considerably since the processor only has to add
- The +1 for negative numbers...
 - makes it so there is only one zero
 - values range from -128 to 127

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

16

2's Complement: 13 and -13

Positive

0 0 0 0 1 1 0 1

Negative

1 1 1 1 0 0 1 1

Add 1

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

17

Just One Zero!

0

0 0 0 0 0 0 0 0


-1

1 1 1 1 1 1 1 1

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

18




Adding Binary Integers

$$1 + 1 = 10$$

Adding Binary Integers

- Computer's add binary numbers the same way that we do with decimal
- Columns are aligned, added, and "1's" are carried to the next column
- In computer processors, this component is called an *adder*



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

20

Adding Base 10 Numbers

	1	1		
	2	7	8	1
+	3	7	2	1
<hr/>				
	6	5	0	2

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

21

Adding Binary Example

182		1	1		1	1	
	1	0	1	1	0	1	1
51	+	0	0	1	1	0	0
<hr/>							
233		1	1	1	0	1	0

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

22


Adding 2's Complement

91		1	1	1	1		1	1
	0	1	0	1	1	0	1	1
-13	+	1	1	1	1	0	0	1
<hr/>								
78		0	1	0	0	1	1	0

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

23



Extending Bytes

Converting from 8-bit to 16-bit and more

Extending Unsigned Integers

- Often in programs, data needs to be moved to an integer with a larger number of bits
- For example, an 8-bit number is moved to a 16-bit representation



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

25

Extending Unsigned Integers

- For unsigned numbers is fairly easy – just add zeros to the left of the number
- This, naturally, is how our number system works anyway: $000456 = 456$



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

26

Unsigned 13 Extended

0 0 0 0 1 1 0 1



0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

27

Extending Signed Integers

- When the data is stored in a signed integer, the conversion is a little more complex
- Simply adding zeroes to the left, will *convert a negative value to a positive one*
- Each type of signed representation has its own set of rules

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

28

2's Complement Extended Incorrectly

-13

1 1 1 1 0 0 1 1

243



0 0 0 0 0 0 0 0 0 0 1 1 1 1 0 0 1 1

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

29

Sign Magnitude Extension

- In signed magnitude, the most-significant bit (msb) stores the negative sign
- The new sign-bit needs to have this value
- Rules:
 - copy the old sign-bit to the new sign-bit
 - fill in the rest of the new bits with zeroes – *including the old sign bit*

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

30

Sign Magnitude Extended: +77

0 1 0 0 1 1 0 1



0 0 0 0 0 0 0 0 0 1 0 0 1 1 0 1

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

31

Sign Magnitude Extended: -77

1 1 0 0 1 1 0 1



1 0 0 0 0 0 0 0 0 1 0 0 1 1 0 1

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

32

2's Complement Extension

- 2's Complement is very simple to convert to a larger representation
- Remember that we inverted the bits and added 1 to get a negative value
- Rule: copy the old most-significant bit to all the new bits

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

33

2's Complement Extended: +77

0 1 0 0 1 1 0 1



0 0 0 0 0 0 0 0 0 1 0 0 1 1 0 1

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

34

2's Complement Extended: -77

1 0 1 1 0 0 1 1



1 1 1 1 1 1 1 1 1 0 1 1 0 0 1 1

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

35




Multiplying
Binary
Numbers

$$11 \times 11 = 1001$$

Multiplying Binary Numbers

- Many processors today provide complex mathematical instructions
- However, the processor only needs to know how to add
- Historically, multiplication was performed with successive additions



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

37

Multiplying Scenario

- Let's say we have two variables: A and B
- Both contain integers that we need to multiply
- Our processor can only add (and subtract using 2's complement)
- How do we multiply the values?

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

38

Multiplying: The Bad Way



- One way of multiplying the values is to create a For Loop using one of the variables – A or B
- Then, inside the loop, continuously add the other variable to a running total

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

39

Multiplying: The Bad Way

```
total = 0;
for (i = 0; i < A; i++)
{
    total += B;
}
```

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

40

Multiplying: The Bad Way

- If one of the operands – A or B – is large, then the computation could take a long time
- This is incredibly inefficient
- Also, given that A and B could contain drastically different values – the number of iterations would vary
- Required time is not constant



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

41

Multiplying: The Best Way



- Computers can perform multiplication using long multiplication – just like you do
- The number of additions is then fixed to 8, 16, 32, 64 depending on the size of the integer
- The following example multiplies 2 unsigned 4-bit numbers

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

42

Unsigned Integer: 13×10

$$\begin{array}{r} 1101 \\ \times 1010 \\ \hline \end{array}$$

$$\begin{array}{r} + \\ \hline \end{array}$$

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

43

Unsigned Integer: 13×10

$$\begin{array}{r} 1101 \\ \times 1010 \\ \hline \end{array}$$

$$\begin{array}{r} + \\ \hline \end{array}$$

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

44

Unsigned Integer: 13×10

$$\begin{array}{r} 1101 \\ \times 1010 \\ \hline 0000 \\ \hline \end{array}$$

$$\begin{array}{r} + \\ \hline \end{array}$$

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

45

Unsigned Integer: 13×10

$$\begin{array}{r} 1101 \\ \times 1010 \\ \hline 0000 \\ \hline \end{array}$$

$$\begin{array}{r} + \\ \hline \end{array}$$

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

46

Unsigned Integer: 13×10

$$\begin{array}{r} 1101 \\ \times 1010 \\ \hline 0000 \\ 1101 \\ \hline \end{array}$$

$$\begin{array}{r} + \\ \hline \end{array}$$

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

47

Unsigned Integer: 13×10

$$\begin{array}{r} 1101 \\ \times 1010 \\ \hline 0000 \\ 1101 \\ \hline \end{array}$$

$$\begin{array}{r} + \\ \hline \end{array}$$

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

48

Unsigned Integer: 13×10

$$\begin{array}{r}
 1101 \\
 \times 1010 \\
 \hline
 0000 \\
 1101 \\
 0000 \\
 \hline
 \end{array}$$

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

49

Unsigned Integer: 13×10

$$\begin{array}{r}
 1101 \\
 \times 1010 \\
 \hline
 0000 \\
 1101 \\
 0000 \\
 \hline
 \end{array}$$

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

50

Unsigned Integer: 13×10

$$\begin{array}{r}
 1101 \\
 \times 1010 \\
 \hline
 0000 \\
 1101 \\
 0000 \\
 \hline
 1101
 \end{array}$$

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

51

Unsigned Integer: 13×10

$$\begin{array}{r}
 1101 \\
 \times 1010 \\
 \hline
 0000 \\
 1101 \\
 0000 \\
 \hline
 1101 \\
 \hline
 10000010
 \end{array}$$

130

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

52

Multiplication Doubles the Bit-Count

- When two numbers are multiplied, the product will have twice the number of digits
- Examples:
 - 8-bit \times 8-bit \rightarrow 16-bit
 - 16-bit \times 16-bit \rightarrow 32-bit
- Often processors...
 - will store the result in the original bit-size
 - and flag an overflow if it does not fit

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

53



x86
Mathematics

Complex Math is Complex

Add & Subtract

- The Add and Subtract instructions take two operands and store the result in the second operand
- This is the same as the **+=** and **-=** operators used in Visual Basic .NET, C, C++, Java, etc...



6/7/2019

Sacramento State - Cook - CSIS 35 - Spring 2019

55

Addition

Immediate, Register, Memory

ADD *value*, *target*

Register, Memory

6/7/2019

Sacramento State - Cook - CSIS 35 - Spring 2019

56

Subtraction

Immediate, Register, Memory

SUB *value*, *target*

Register, Memory

6/7/2019

Sacramento State - Cook - CSIS 35 - Spring 2019

57

Negate (2's complement)

NEG *register*

6/7/2019

Sacramento State - Cook - CSIS 35 - Spring 2019

58

Example: Simple Add

```
MOV $17, %rax
ADD $2, %rax
```

Move value into RAX

RAX += 2


6/7/2019

Sacramento State - Cook - CSIS 35 - Spring 2019

59

Multiplication & Division

- Multiplication and division are far more complex
- The x86 treats this type of math differently than add/subtract
- It requires specific registers to be used: *RAX*, *RDX*



6/7/2019

Sacramento State - Cook - CSIS 35 - Spring 2019

60

Multiplication Review

- Remember: when two n bit numbers are multiplied, result will be $2n$ bits
- So...
 - two 8-bit numbers \rightarrow 16-bit
 - two 16-bit numbers \rightarrow 32-bit
 - two 32-bit numbers \rightarrow 64-bit
 - two 64-bit numbers \rightarrow 128-bit



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

61

Multiplication on the x86


- Instruction inputs are strange
 - first operand is must be stored in **RAX**
 - second operand must be a register (can't be a immediate)
- Result is stored into two registers
 - rax** will contain the lower 8 bytes
 - rdx** will contain the upper 8 bytes

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

62

x86 Multiplication



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

63

Multiply - Signed

IMUL *operand*

Register, Memory

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

64

Multiply - Unsigned

MUL *operand*

Register, Memory

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

65

Unsigned Multiply: 1846 by 42

```
MOV    $1846, %rax    #First operand
MOV    $42, %rbx      #Need register for MUL
MUL    %rbx           #RAX gets low 8 bytes
                        #RDX gets high 8 bytes
```

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

66

Multiplication Tips

- Even though you are just using RAX as input, **both** RAX and RDX will change
- Be aware that you might lose important data, and backup to memory if needed



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

67

Additional x86 Multiply Instructions

- x86 also contains versions of the IMUL instruction that take multiple operands
- Allows "short" multiplication – just stored in 1 register
- Please note: these do **not** exist for MUL



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

68

IMUL (few more combos)

```
IMUL immediate, reg
IMUL memory, reg
IMUL reg, reg
```

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

69

Signed Multiply: 1846 by 42

```
MOV    $1846, %rax
IMUL   $42, %rax
```

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

70

Division on the x86

- Division on the x86 is very interesting
- Like multiplication, it uses 2 registers
- The dividend (number being divided) uses **two** registers
 - RAX contains the lower 8 bytes
 - RDX contains the upper 8 bytes




6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

71

Division on the x86

- These two registers are used for the result
- The output contains:
 - RAX will contain the quotient
 - RDX will contain the remainder

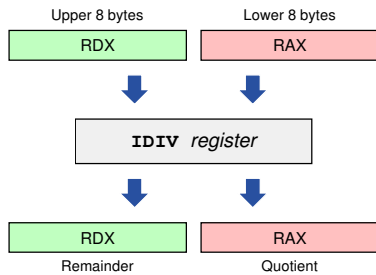


6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

72

x86 Division



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

73

Divide - Signed

IDIV *denominator*

Register, Memory

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

74

Divide - Unsigned

DIV *denominator*

Register, Memory

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

75

Dividing Rules

- Sign Magnitude
 - create a new sign that is a XOR of the old ones
 - clear the old sign bits and expand with zeros
- 2's complement
 - numerator **must** be expanded to the destination size (twice the original)
 - this must be done **beforehand**
 - otherwise the result will be incorrect

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

76

Division Tips

- Even you don't store anything into RDX, you **must** clear it
- For signed-division...
 - RDX needs to contain *the sign-extended* value of RAX
 - this is covered in the next section



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

77

CWD (16 bit): Extend AX → DX

CWD

Convert Word Double

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

78

CDQ (32 bit): Extend EAX → EDX

CDQ

Convert Double Quad

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

79

CQO (64 bit): Extend RAX → RDX

CQO

Convert Quad Oct

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

80

Divide 16-bit: -1846 by 42

```
MOV    $-1846, %ax    #AX is the dividend
MOV    $42, %bx       #Divisor
CWD    #Extend the sign to DX
IDIV   %bx            #AX gets quotient
                        #DX gets remainder
```

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

81

Divide 64-bit: -1846 by 42

```
MOV    $-1846, %rax   #RAX is the dividend
MOV    $42, %rbx      #Divisor
CQO    #Extend the sign to RDX
IDIV   %rbx           #RAX gets quotient
                        #RDX gets remainder
```

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

82



Numbers in Programs

The primitive types are pretty primitive

Primitive Data Types

- Most popular program languages hide the true nature of the computer from you
- However, most of the language's primitive data types are the same types recognized by the processor



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

84

Integer Data Types

- Integer data types are stored in simple binary numbers
- The number of bytes used varies: 1, 2, 4, etc....
- Languages often have a unique name for each – short, int, long, etc...

1234

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

85

Floating-Point Data Type

- Floating-point numbers are usually stored using the IEEE 754 standard
- Languages often have unique names for them such as float, double, real

3.14
1.618
2.71

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

86

Floating-Point Data Type


- This is not always the case
 - some languages implement their own structures
 - e.g. COBOL
- Why?
 - some processors do not have floating-point instructions
 - or the language needs more precision and control

3.14
1.618
2.71

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

87




Floating Point Numbers

Real numbers are *real* complex

Floating Point Numbers

- Often, programs need to perform mathematics on *real* numbers
- Floating point numbers* are used to represent quantities that cannot be represented by integers




6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

89

Floating Point Numbers

- Why?
 - regular binary numbers can only store whole positive and negative values
 - many numbers outside the range representable within the system's bit width (too large/small)



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

90

IEEE 754

- Practically modern computers use the *IEEE 754 Standard* to store floating-point numbers
- Represent by a mantissa and an exponent
 - similar to scientific notation
 - the value of a number is: $\text{mantissa} \times 2^{\text{exponent}}$
 - uses signed magnitude

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

91

IEEE 754


- Comes in three forms:
 - single-precision: 32-bit
 - double-precision: 64-bit
 - quad-precision: 128-bit
- Also supports special values:
 - negative and positive *infinity*
 - and "not a number" for errors (e.g. 1/0)

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

92

IEEE 754 Single Precision (32 bit)



6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

93

Interpretation: Invalid Numbers

$NaN \rightarrow 1/0$

$Naan \rightarrow$ 

6/7/2019

Sacramento State - Cook - CSc 35 - Spring 2019

94