




Part 7

Subroutines & Operating Systems




The System Stack

Pile of... Data

The System Stack

- The processor maintains a stack in memory
- It is used to allow *subroutines* which are similar to the function you are used to programming in Java



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

3

Implementing in Memory

- On a processor, the stack stores integers
 - size of the integer the bit-size of the system
 - 64-bit system → 64-bit integer
- A fixed location pointer (S0) defines the bottom of the stack
- A *stack pointer* (SP) gives the location of the top of the stack

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

4

Approaches

- Growing upwards
 - Bottom Pointer (S0) points to the *lowest* address in the stack buffer
 - stack grows towards *higher* addresses
- Grow downwards
 - Bottom Pointer (S0) points to the *highest* address in the stack buffer
 - stack grows towards *lower* addresses

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

5

Size of the Stack


- As an abstract data structure...
 - stacks are assumed to be infinitely deep
 - so, an arbitrary amount of data can stored
- However...
 - stacks are implemented using memory buffers
 - which are finite in size
- If the amount of data *exceeds* the allocated space, a *stack overflow* error occurs

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

6


Subroutine Call Basics



Organizing Your Program

Subroutine Call

- The stack is essential for subroutines to work
- How?
 - used to save the return addresses for call instructions
 - backup and restore registers
 - pass data between subroutines (*we won't get to this*)



6/26/2019 Sacramento State - Cook - CSc 35 - Summer 2019 8

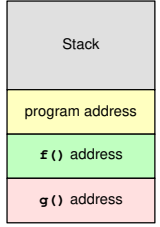
When you call a subroutine...

1. Processor pushes the program counter (PC) – an address – on the stack
2. PC is set to the address of the subroutine
3. Subroutine executes and ends with a "return" instruction
4. Processor pops & restores the original PC
5. Execution continues after the initial call

6/26/2019 Sacramento State - Cook - CSc 35 - Summer 2019 9

Nesting is Possible

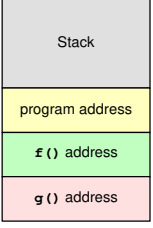
- Subroutines can call other subroutines
- Here: the program calls f() which then calls g(), etc...
- The stack stores the return addresses of the callers
- Just like the "history button" in your web browser, you can store many return addresses



6/26/2019 Sacramento State - Cook - CSc 35 - Summer 2019 10

Nesting is Possible


- Each time a subroutine completes, the processor pops the top of the stack
- ...then returns to the *caller*
- This allows normal function calls and recursion (a powerful tool)



6/26/2019 Sacramento State - Cook - CSc 35 - Summer 2019 11

Passing Parameters

- Useful subroutines...
 - need you to be able to pass data into them
 - and be able to read data from it
- One of the easiest ways to accomplish this is by using the processors registers
- *Incredibly efficient!*



6/26/2019 Sacramento State - Cook - CSc 35 - Summer 2019 12

Passing Through Registers

- However... processor might not have enough registers for each parameter
- *Other forms of passing data have to be used*



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

13

x64 Subroutines



Organizing Your Programs ... with Intel

Instruction: Call

- The *Call Instruction* is used to transfer control to a subroutine
- Other processors call it different names such as JSR (Jump Subroutine)
- The stack is used to save the current PC

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

15

Instruction: Call

Usually a label
(which is an address)

CALL *address*

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

16

Instruction: Return

- The Return Instruction is used mark the end of subroutine
- When the instruction is executed...
 - the old program counter is read from the system stack
 - the current program counter is updated – restoring execution after the initial call

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

17

Instruction: Return

- Do not forget this!
- If you do...
 - execution will simply continue, in memory, until a return instruction is encountered
 - often is can run past the end of your program
 - ...and run data!

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

18

Instruction: Return

RET

No arguments!

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

19

Subroutine Example

```
_start:
    mov $4, %rax
    mov $12, %rbx
    call AddIt
    add $1, %rbx
    ...
AddIt:
    add %rax, %rbx
    ret
```

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

20

Saving Registers & Lost Data



Avoiding horrible side-effects

Saving Registers & Lost Data

- Each subroutine will use the registers as it needs
- So, when a sub is called, *it may modify the caller's registers*
- Some processors have few registers – so its *very* likely
- This can lead to hard-to-fix bugs if caution is not used – e.g. loop counter gets changed



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

22

Two Solutions

- Caller saves values
 - caller saves all their registers to memory before making the subroutine call
 - after, it restores the values before continuing
 - not recursion friendly – it pushes all of them!
- Subroutine saves the values
 - push registers (it will change) onto the stack
 - before it returns, it pops (and restores) the old values off the stack

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

23

Saving Registers... How nice! :-)

```
DoSomething:
    push %rax
    push %rbx
    push %rcx
    ...
    pop %rcx
    pop %rbx
    pop %rax
    ret
```

Save registers

Your code

Restore them.
Note the reverse order

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

24




Operating Systems

The master software

What is an operating system?


- The operating system is simply a series of programs
- These programs, however, run with special privileges which are needed by the OS
- Processors support two modes for executing programs



6/26/2019 Sacramento State - Cook - CSc 35 - Summer 2019 26

Execution Modes


- *Privileged (supervisor) mode*
 - can run special instructions
 - can talk to all the hardware
 - etc...
- *User mode*
 - can only execute certain instructions
 - can't talk to all the hardware



6/26/2019 Sacramento State - Cook - CSc 35 - Summer 2019 27

Vector Tables

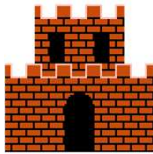
- Programs (and hardware) often need to talk to the operating system
- Examples:
 - software needs talk to the OS
 - USB port notifies the OS that a device was plugged in



6/26/2019 Sacramento State - Cook - CSc 35 - Summer 2019 28


Vector Tables

- But how does this happen?
- The processor can be interrupted – alerted – that something must be handled
- It then runs a special program that handles the event



6/26/2019 Sacramento State - Cook - CSc 35 - Summer 2019 29


Vector Table



- During an interrupt, the device sends the processor an interrupt number
- The processor looks up the number in the vector table
- Table contains the address of Interrupt Service Routine (ISR) to execute

6/26/2019 Sacramento State - Cook - CSc 35 - Summer 2019 30

How It Works



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

31

The Processor Actions



1. Backup the register file
2. Execute *Interrupt Service Routine (ISR)*
3. Once completed: restore the original executing program & register file

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

32

The Kernel

- All these Interrupt Service Routines belong to the *kernel* – the core of the operating system
- Vast majority of the operating system is hidden from the end user



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

33



Interact with Applications

How do WE talk to the OS

Interact with Applications

- Software also needs to talk to the operating system
- For example:
 - draw a button
 - print a document
 - close this program
 - etc...



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

35

Interact with Applications

- Software can interrupt itself with a specific number
- This interrupt is *designated specifically for software*
- The operating system then handles the software's request



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

36

Application Program Interface

- Programs "talk" to the OS using Application Program Interface (API)
- Benefits:
 - makes applications faster and smaller
 - also makes the system more secure since apps do not directly talk to IO
 - Application → Operating System → IO

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

37

The Changing Interrupts...

- In the 32-bit Intel x86 the program can invoke any interrupt number
- ... which turned out not to be particularly safe
- So, the Intel x64 hardcoded the interrupt number



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

38

Instruction: syscall (64-bit)

SYSCALL

Calls interrupt number reserved for programs needing attention

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

39

Returning from an Interrupt

IRET

Interrupt Return
(You will never use this unless you write an operating system)

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

40


Subroutine vs. Interrupt

Subroutine	Interrupt
Executes code	Executes code
Returns when complete	Returns when complete
Called by the application	Executed by the processor
Part of the application	Handles events for the OS

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

41




Linux System Calls

How software and hardware "talk"

Interrupts on the Linux

- Linux, like other operating systems communicate with applications using *interrupts*
- Applications do not know where (in memory) to contact the kernel – so they ask the processor to do it




6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

43

How It Works

1. Fill the registers
2. Interrupt using syscall (or int 0x80 if on 32-bit)
3. Any results will be stored in the registers



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

44

How to Call Linux – 64 bit



- The *rax* register must contain *system call number*
- This number indicates what you asking the OS to do
- There are only 329 total calls in the entire 64-bit UNIX operating system!

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

45

How to Call Linux – 64 bit



- Different registers are used to hold data
- The order is also quite odd: *rdi, rsi, rdx, r10, r8*


6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

46

Kernels are Simple!

- Linux only has **1** write and **1** read system call
- The location, number of bytes, and device only change
"write x many bytes from address y to device z"
- So, writing to the screen, a file, a port, etc...use the *same* call!



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

47

Some Linux 64 Calls

System Call	rax	rdi	rsi	rdx
read	0	fd (device)	address	max bytes
write	1	fd (device)	address	count
open	2	address	flags	mode
close	3	fd (device)		
get pid	39			
exit	60	error code		

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

48

Linux 64: Sys Write

```
mov $1, %rax
mov $1, %rdi
mov $address, %rsi
mov $length, %rdx
syscall
```

Linux command for WRITE

1 = Screen

Call Linux

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

49

Linux 64: Sys Read

```
mov $0, %rax
mov $0, %rdi
mov $address, %rsi
mov $maxBytes, %rdx
syscall
```

Linux command for READ

0 = Keyboard

Maximum number of bytes to read

Call Linux

6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

50

Null-Terminated Strings

- Notice that UNIX requires that you specify the number of bytes before you write
- So, you need to count the bytes every time




6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

51

Null-Terminated Strings

- That's why C-style null-terminated strings are used
- The programming language (or library) counts the bytes for you!



6/26/2019

Sacramento State - Cook - CSc 35 - Summer 2019

52